

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
!ls "/content/drive/My Drive"
```

```
'Colab Notebooks' 'Fixed Resume.pdf' 'M.Tech DA'
```

```
data=pd.read_csv('/content/drive/My Drive/Colab
Notebooks/Assignment3.csv')
```

```
data
```

```
{
  "summary": {
    "name": "data",
    "rows": 101,
    "fields": [
      {
        "column": "x1",
        "properties": {
          "dtype": "number",
          "std": 0.3801148862287207,
          "min": 6.68,
          "max": 8.37,
          "num_unique_values": 71,
          "samples": [
            7.0,
            7.11,
            8.09
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "x2",
        "properties": {
          "dtype": "number",
          "std": 292.8501773932321,
          "min": -466.86,
          "max": 546.88,
          "num_unique_values": 101,
          "samples": [
            361.89,
            73.03,
            193.86
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "x3",
        "properties": {
          "dtype": "number",
          "std": 55.81221280382463,
          "min": 9.8,
          "max": 195.81,
          "num_unique_values": 82,
          "samples": [
            168.33,
            135.66,
            87.58
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "x4",
        "properties": {
          "dtype": "number",
          "std": 4.942089001927251,
          "min": 86.83,
          "max": 108.85,
          "num_unique_values": 100,
          "samples": [
            107.59,
            95.67,
            104.6
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "x5",
        "properties": {
          "dtype": "number",
          "std": 30.557704102239096,
          "min": 0.0,
          "max": 100.0,
          "num_unique_values":

```

```

51,\n          \"samples\": [\n          1.44,\n          4.84,\n0.64\n          ],\n          \"semantic_type\": \"\",\n\"description\": \"\"\n          }\n          },\n          {\n          \"column\":\n\"y\",\n          \"properties\": {\n          \"dtype\": \"number\",\n\"std\": 1022.7661225514418,\n          \"min\": 8062.54,\n\"max\": 12631.05,\n          \"num_unique_values\": 101,\n\"samples\": [\n          12266.88,\n          10543.83,\n11065.37\n          ],\n          \"semantic_type\": \"\",\n\"description\": \"\"\n          }\n          }\n          ]\n          }\", \"type\": \"dataframe\", \"variable_name\": \"data\"}

```

Statistics of the data

```
data.describe()
```

```

{"summary": "{\n  \"name\": \"data\",\n  \"rows\": 8,\n  \"fields\": [\n    {\n      \"column\": \"x1\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 33.50474167750145,\n        \"min\": 0.3801148862287207,\n        \"max\": 101.0,\n        \"num_unique_values\": 8,\n        \"samples\": [\n          7.54871287128713,\n          7.53,\n          101.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"x2\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 312.6476261638648,\n        \"min\": -466.86,\n        \"max\": 546.88,\n        \"num_unique_values\": 8,\n        \"samples\": [\n          23.755049504950488,\n          38.95,\n          101.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"x3\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 59.96719780751142,\n        \"min\": 9.8,\n        \"max\": 195.81,\n        \"num_unique_values\": 8,\n        \"samples\": [\n          111.37138613861389,\n          104.18,\n          101.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"x4\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 33.61372219855231,\n        \"min\": 4.942089001927251,\n        \"max\": 108.85,\n        \"num_unique_values\": 8,\n        \"samples\": [\n          98.13376237623763,\n          97.9,\n          101.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"x5\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 38.79923873735032,\n        \"min\": 0.0,\n        \"max\": 101.0,\n        \"num_unique_values\": 8,\n        \"samples\": [\n          101.0,\n          34.0,\n          25.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ]\n},\n  {\n    \"column\": \"y\",\n    \"properties\": {\n      \"dtype\": \"number\",\n      \"std\": 4667.185911152322,\n      \"min\": 101.0,\n      \"max\": 12631.05,\n      \"num_unique_values\": 8,\n      \"samples\": [\n

```

```
10244.460297029704,\n          10187.66,\n          101.0\n ],\n      \"semantic_type\": \"\",\n      \"description\": \"\"\n }\n ]\n }\", \"type\": \"dataframe\"}
```


Check for any null values

```
data.isna().sum()
```

```
x1      0\nx2      0\nx3      0\nx4      0\nx5      0\ny      0\ndtype: int64
```

Visualization of the relationship of the features using pairplot

```
import seaborn as sns\nsns.pairplot(data)\n<seaborn.axisgrid.PairGrid at 0x7d670440ddb0>
```



### Observations 1:

1. As we can observe from the pairplot that only features x1 and features x4 are having linear relationship with dependent features.
2. Feature x2, feature x3 and feature x5 not have linear relationship with the dependent features.

## Task 1

Fit OLS on the data directly and evaluate the baseline SSE loss. You will observe that the loss is very high, but that's ok. You will strive hard to apply creative ways to reduce the loss.

```
# separating the dependent and independent features
y=data['y']      # independent features
X=data.drop(columns=['y']) # independent features
```

```
y      # dependent feature
```

```
0      9131.40
1      9001.86
2      8595.85
3      9469.94
4      9448.98
```

```
...
96     11168.68
97     12605.81
98     12467.96
99     12631.05
100    10327.89
```

```
Name: y, Length: 101, dtype: float64
```

```
X      # independent features
```

```
{"summary":{"\n  \"name\": \"X\",\n  \"rows\": 101,\n  \"fields\": [\n    {\n      \"column\": \"x1\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.3801148862287207,\n        \"min\": 6.68,\n        \"max\": 8.37,\n        \"num_unique_values\": 71,\n        \"samples\": [\n          7.0,\n          7.11,\n          8.09\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"x2\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 292.8501773932321,\n        \"min\": -466.86,\n        \"max\": 546.88,\n        \"num_unique_values\": 101,\n        \"samples\": [\n          361.89,\n          73.03,\n          193.86\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"x3\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 55.81221280382463,\n        \"min\": 9.8,\n        \"max\": 195.81,\n        \"num_unique_values\": 82,\n        \"samples\": [\n          168.33,\n          135.66,\n          87.58\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"x4\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 4.942089001927251,\n        \"min\": 86.83,\n        \"max\": 108.85,\n        \"num_unique_values\": 100,\n        \"samples\": [\n          107.59,\n          95.67,\n          104.6\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"x5\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 30.557704102239096,\n        \"min\": 0.0,\n        \"max\": 100.0,\n        \"num_unique_values\": 51,\n        \"samples\": [\n          1.44,\n          4.84,\n          0.64\n        ],\n        \"semantic_type\": \"\",
```

```
\n"description\\": \\\"\\\"\\n      }\\n      }\\n  ]\\n}\\", "type": "dataframe", "variable_name": "X"}\n\nimport statsmodels.api as sm
```

Model Summary

```
X=sm.add_constant(X)\nmodel=sm.OLS(y,X).fit()\nprint(model.summary())
```

OLS Regression Results						
=====						
=====						
Dep. Variable:	y	R-squared:				
0.999						
Model:	OLS	Adj. R-squared:				
0.999						
Method:	Least Squares	F-statistic:				
2.763e+04						
Date:	Sun, 25 Aug 2024	Prob (F-statistic):				
1.47e-148						
Time:	10:15:01	Log-Likelihood:				
-474.98						
No. Observations:	101	AIC:				
962.0						
Df Residuals:	95	BIC:				
977.6						
Df Model:	5					
Covariance Type:	nonrobust					
=====						
=====						
	coef	std err	t	P> t	[0.025	
0.975]						
-----						
-----						
const	-9655.3103	83.303	-115.906	0.000	-9820.688	-
9489.933						
x1	-1067.3690	1147.895	-0.930	0.355	-3346.229	
1211.491						
x2	0.1007	0.014	7.289	0.000	0.073	
0.128						
x3	-0.0572	0.053	-1.083	0.282	-0.162	
0.048						
x4	284.3633	88.453	3.215	0.002	108.763	
459.964						
x5	1.6285	0.090	18.017	0.000	1.449	

```

1.808
=====
=====
Omnibus:                24.930    Durbin-Watson:
1.535
Prob(Omnibus):          0.000    Jarque-Bera (JB):
36.395
Skew:                   1.149    Prob(JB):
1.25e-08
Kurtosis:               4.836    Cond. No.
1.23e+05
=====
=====

```

#### Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.23e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Here R squared and adjusted R squared are same.

Can we have the same Rsquared and AdjRsquared ?

- Yes, it can be same in two case:
  1. Single predictor: When there is only one independent variable. the R squared and adjusted R square will be the same because the adjustment factor for k becomes negligible.
  2. Perfect fit: If the model perfectly predicts the dependent variable (R squared=1), the R squared and adjusted R squared will be same.

Sum of squared Errors(SSE) and Root Mean Squared Error(RMSE)

```

predictions=model.predict(X)
residuals=y-predictions
## SSE
sse1=np.sum(residuals**2)
rmse1=np.sqrt(sse1/101)
print("Sum of Squared Errors(SSE):",sse1)
print("Root Mean Squared Error(RMSE):",rmse1)

Sum of Squared Errors(SSE): 71877.84134016861
Root Mean Squared Error(RMSE): 26.67698999975848

```

## Task 2

Perform EDA on the dataset to understand the predictor features and how are they influencing each other. Also, study how each individual predictor influence the output variable. You may use correlation study to estimate the influence. Add necessary visualization and its representative interpretations to substantiate your inferences. The outcome of this step is figure out the requires features and their respective transformation.

Statistics of the features and outputs

```
# description of the data in a DataFrame. The description includes
summary statistics for each numerical column in the DataFrame,
# such as the number of non-empty values, the mean, the standard
deviation, the minimum and maximum values, and the percentiles
data.describe()

{"summary":{"\n  \"name\": \"data\",\n  \"rows\": 8,\n  \"fields\": [\n    {\n      \"column\": \"x1\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 33.50474167750145,\n        \"min\": 0.3801148862287207,\n        \"max\": 101.0,\n        \"num_unique_values\": 8,\n        \"samples\": [\n          7.54871287128713,\n          7.53,\n          101.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"x2\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 312.6476261638648,\n        \"min\": -466.86,\n        \"max\": 546.88,\n        \"num_unique_values\": 8,\n        \"samples\": [\n          23.755049504950488,\n          38.95,\n          101.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"x3\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 59.96719780751142,\n        \"min\": 9.8,\n        \"max\": 195.81,\n        \"num_unique_values\": 8,\n        \"samples\": [\n          111.37138613861389,\n          104.18,\n          101.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"x4\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 33.61372219855231,\n        \"min\": 4.942089001927251,\n        \"max\": 108.85,\n        \"num_unique_values\": 8,\n        \"samples\": [\n          98.13376237623763,\n          97.9,\n          101.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"x5\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 38.79923873735032,\n        \"min\": 0.0,\n        \"max\": 101.0,\n        \"num_unique_values\": 8,\n        \"samples\": [\n          101.0,\n          34.0,\n          25.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ]\n},\n  \"column\": \"y\",\n  \"properties\": {\n    \"dtype\": \"number\",\n    \"std\": 4667.185911152322,\n    \"min\": 101.0,\n    \"max\": \"\"
```



```
12631.05,\n          \"num_unique_values\": 8,\n          \"samples\": [\n10244.460297029704,\n          10187.66,\n          101.0\n],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n}\n ]\n}","type":"dataframe"}
```

Check for any null values


```
data.isna().sum()
```

```
x1      0
x2      0
x3      0
x4      0
x5      0
y        0
dtype: int64
```

From above two(.describe() and .isna().sum()) observations we can see that there is no null values in the dataset

Bivariate Analysis Using Correlation Matrix

```
correlation_matrix = data.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```



#### Observations 2:

1. Features 1 and feature 4 is highly positive correlated , so we can drop one feature among feautre 1 and feature 4.
2. Features 3 and feature 5 is highly positive correlated , so we can drop one feature among feautre 3 and feature 5.
3. Independent feature1(x1) and dependent feature y, independent feature 4 and dependent feature y is highly positive correlated.
4. Independent feature(x2 and x5) and y is positively correlated.
5. Independent feature(x3) and dependent feature y is negatively correlated.

```
# Variance Inflation Factor
from statsmodels.stats.outliers_influence import
variance_inflation_factor
import statsmodels.api as sm

vif=pd.DataFrame()
vif['Variables']=X.columns
vif['vif']=[variance_inflation_factor(X.values,i) for i in
range(X.shape[1])]

vif
```

```
{
  "summary": {
    "name": "vif",
    "rows": 6,
    "fields": [
      {
        "column": "Variables",
        "properties": {
          "dtype": "string",
          "num_unique_values": 6,
          "samples": [
            "const",
            "x1",
            "x5"
          ],
          "semantic_type": ""
        },
        "description": ""
      },
      {
        "column": "vif",
        "properties": {
          "dtype": "number",
          "std": 12903.08705860959,
          "min": 1.0082661211823698,
          "max": 25256.36370857026,
          "num_unique_values": 6,
          "samples": [
            926.3399082382757,
            25163.014234909457,
            1.0082661211823698
          ],
          "semantic_type": "",
          "description": ""
        }
      }
    ]
  },
  "type": "dataframe",
  "variable_name": "vif"
}
```

Variance Inflation Factor(VIF):

- It measures how much the variance of a regression coefficient is inflated due to multicollinearity with other variables.
- $VIF = 1/(1-R^2)$ , where  $R^2$  is used to calculate the accuracy of the model.
- $VIF = 1$  No multicollinearity (the variable is not correlated with other predictors)
- $1 < VIF < 5$ : Moderate correlation.
- $VIF \geq 5$  Indicates problematic multicollinearity.
- $VIF \geq 10$  Strong multicollinearity, seen as a serious issue.

### Observations 3:

1. We can see that variables x1 and variable x4 are highly correlated. The same observation we have seen using the correlation matrix.
2. We will remove the features const and feature x1.

```
X = X.drop(columns=['x1', 'const'])
X
```

```
{
  "summary": {
    "name": "X",
    "rows": 101,
    "fields": [
      {
        "column": "x2",
        "properties": {
          "dtype": "number",
          "std": 292.8501773932321,
          "min": -466.86,
          "max": 546.88,
          "num_unique_values": 101,
          "samples": [
            361.89,
            73.03,
            193.86
          ],
          "semantic_type": "",
          "description": ""
        }
      },
      {
        "column": "x3",
        "properties": {
          "dtype": "number",
          "std": 55.81221280382463,
          "min": 9.8,
          "max": 195.81,
          "num_unique_values": 82,
          "samples": [
            168.33,
            135.66,
            87.58
          ],
          "semantic_type": "",
          "description": ""
        }
      },
      {
        "column": "x4",
        "properties": {
          "dtype": "number",
          "std": 4.942089001927251,
          "min": 86.83,
          "max": 108.85,
          "num_unique_values": 82,
          "samples": [
            168.33,
            135.66,
            87.58
          ],
          "semantic_type": "",
          "description": ""
        }
      }
    ]
  }
}
```

```

\ "num_unique_values\ ": 100,\n          \ "samples\ ": [\n
107.59,\n          95.67,\n          104.6\n          ],\n
\ "semantic_type\ ": \ "\",\n          \ "description\ ": \ "\",\n          }\n
n          },\n          {\n          \ "column\ ": \ "x5\ ",\n          \ "properties\ ": {\n
\ "dtype\ ": \ "number\ ",\n          \ "std\ ": 30.557704102239096,\n
\ "min\ ": 0.0,\n          \ "max\ ": 100.0,\n          \ "num_unique_values\ ":
51,\n          \ "samples\ ": [\n          1.44,\n          4.84,\n
0.64\n          ],\n          \ "semantic_type\ ": \ "\",\n
\ "description\ ": \ "\",\n          }\n          }\n          ]\n
n}", "type": "dataframe", "variable_name": "X"}

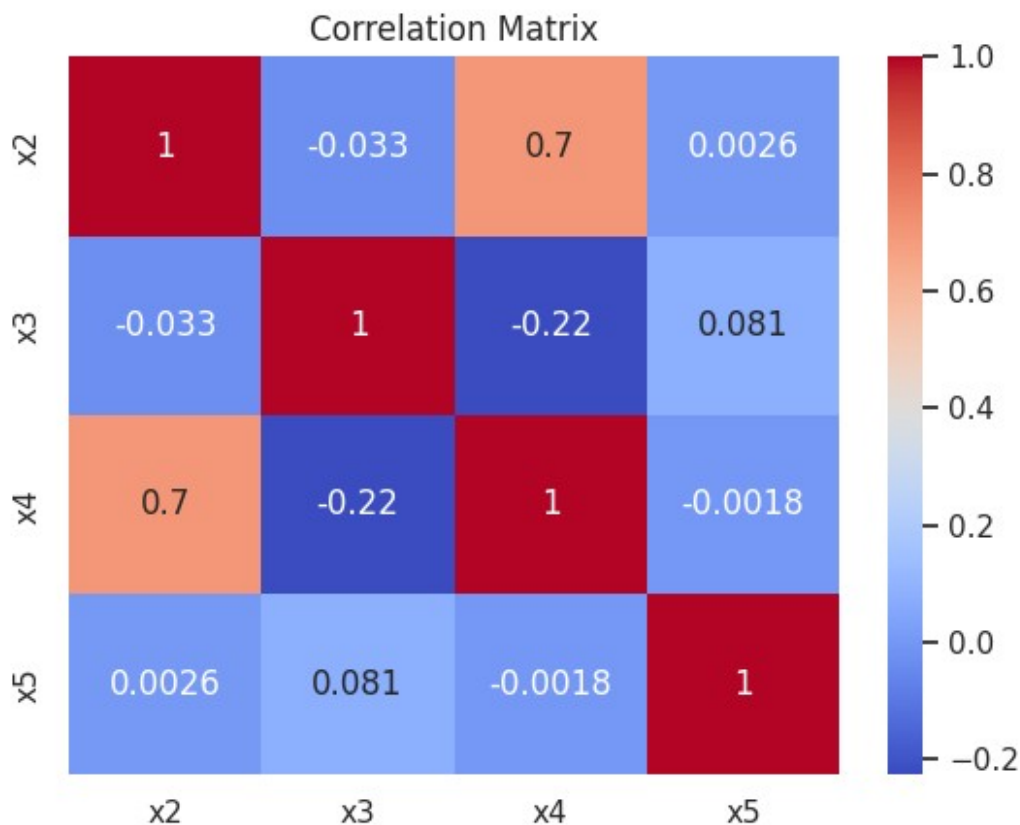
```

Let's see the Correlation Matrix after removing the features x1 from the dataset.

```

correlation_matrix = X.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()

```



Variance Inflation factor

```

# Variance Inflation Factor
from statsmodels.stats.outliers_influence import
variance_inflation_factor

```

```

import statsmodels.api as sm

vif=pd.DataFrame()
vif['Variables']=X.columns
vif['vif']=[variance_inflation_factor(X.values,i) for i in
range(X.shape[1])]

vif

{"summary":{"\n  \"name\": \"vif\", \n  \"rows\": 4, \n  \"fields\": [\n    \n    \"column\": \"Variables\", \n    \"properties\": {\n      \"dtype\": \"string\", \n      \"num_unique_values\": 4, \n      \"samples\": [\n        \"x3\", \n        \"x5\", \n        \"x2\" \n      ], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n    }, \n    { \n      \"column\": \"vif\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 2.1975557653176843, \n        \"min\": 1.024491999281853, \n        \"max\": 5.740196624898975, \n        \"num_unique_values\": 4, \n        \"samples\": [\n          4.8496431595827945, \n          2.265825309122567, \n          1.024491999281853 \n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n      } \n    } \n  ] \n}, \"type\": \"dataframe\", \"variable_name\": \"vif\"}

```

#### Observations 4:

1. We can observe from **correlation matrix** and **variance inflation factor** that there is **no independent features** which are **highly correlated**.

#### Conclusion of EDA

1. From **observation 1** we have observed that feature x2 , feature x3 and feature x5 not have any linear relationship with the dependent features.
2. So, we add some feature which are having polynomial degree(degree=2) consist of feature x2 ,feature x3 and feature x5.

Adding some extra features

```

X['x6']=X['x2']*X['x5']
X['x7']=X['x2']*X['x4']
X['x8']=X['x2']*X['x3']
X['x9']=X['x3']*X['x4']
X['x10']=X['x3']*X['x5']
X['x11']=X['x4']*X['x5']
X['x12']=X['x2']*X['x2']
X['x13']=X['x3']*X['x3']
X['x14']=X['x5']*X['x5']
X

{"summary":{"\n  \"name\": \"X\", \n  \"rows\": 101, \n  \"fields\": [\n    \n    \"column\": \"x2\", \n    \"properties\": {\n

```

```

\"dtype\": \"number\", \n          \"std\": 292.8501773932321, \n
\"min\": -466.86, \n          \"max\": 546.88, \n
\"num_unique_values\": 101, \n          \"samples\": [ \n
361.89, \n          73.03, \n          193.86 \n          ], \n
\"semantic_type\": \"\", \n          \"description\": \"\" \n          } \n
}, \n          { \n          \"column\": \"x3\", \n          \"properties\": { \n
\"dtype\": \"number\", \n          \"std\": 55.81221280382463, \n
\"min\": 9.8, \n          \"max\": 195.81, \n
\"num_unique_values\": 82, \n          \"samples\": [ \n          168.33, \n
135.66, \n          87.58 \n          ], \n
\"semantic_type\": \"\", \n          \"description\": \"\" \n          } \n
}, \n          { \n          \"column\": \"x4\", \n          \"properties\": { \n
\"dtype\": \"number\", \n          \"std\": 4.942089001927251, \n
\"min\": 86.83, \n          \"max\": 108.85, \n
\"num_unique_values\": 100, \n          \"samples\": [ \n
107.59, \n          95.67, \n          104.6 \n          ], \n
\"semantic_type\": \"\", \n          \"description\": \"\" \n          } \n
}, \n          { \n          \"column\": \"x5\", \n          \"properties\": { \n
\"dtype\": \"number\", \n          \"std\": 30.557704102239096, \n
\"min\": 0.0, \n          \"max\": 100.0, \n          \"num_unique_values\": 51, \n
\"samples\": [ \n          1.44, \n          4.84, \n
0.64 \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\" \n          } \n          }, \n          { \n          \"column\": \"x6\", \n          \"properties\": { \n          \"dtype\": \"number\", \n
\"std\": 19555.696239116514, \n          \"min\": -46686.0, \n
\"max\": 54688.0, \n          \"num_unique_values\": 101, \n
\"samples\": [ \n          16733.7936, \n          143.1388, \n
1985.1264 \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\" \n          } \n          }, \n          { \n          \"column\": \"x7\", \n          \"properties\": { \n          \"dtype\": \"number\", \n
\"std\": 28863.650047070958, \n          \"min\": -43149.2355, \n
\"max\": 59516.9504, \n          \"num_unique_values\": 101, \n
\"samples\": [ \n          38935.7451, \n          7295.697, \n
19831.878 \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\" \n          } \n          }, \n          { \n          \"column\": \"x8\", \n          \"properties\": { \n          \"dtype\": \"number\", \n
\"std\": 37949.18885334589, \n          \"min\": -86952.675, \n
\"max\": 96174.59250000001, \n          \"num_unique_values\": 101, \n
\"samples\": [ \n          22151.2869, \n          2850.3609, \n
19791.167400000002 \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\" \n          } \n          }, \n          { \n          \"column\": \"x9\", \n          \"properties\": { \n          \"dtype\": \"number\", \n
\"std\": 5401.921458180213, \n          \"min\": 976.08, \n
\"max\": 20034.4095, \n          \"num_unique_values\": 101, \n
\"samples\": [ \n          6585.5839000000005, \n          3899.097, \n
10443.807 \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\" \n          } \n          }, \n          { \n          \"column\": \"x10\", \n          \"properties\": { \n          \"dtype\": \"number\", \n
\"std\": 4751.194534772334, \n          \"min\": 0.0, \n          \"max\":

```

```

18625.0,\n          \"num_unique_values\": 101,\n          \"samples\": [\n
n          2830.3504000000003,\n          76.4988,\n
1045.4016000000001\n          ],\n          \"semantic_type\": \"\",\n
\"description\": \"\"\n          }\n          },\n          {\n          \"column\":\n
\"x11\",\n          \"properties\": {\n          \"dtype\": \"number\",\n
\"std\": 3003.8945434339703,\n          \"min\": 0.0,\n          \"max\":\n
10883.0,\n          \"num_unique_values\": 101,\n          \"samples\": [\n
n          4974.9616000000005,\n          195.804,\n
1047.552\n          ],\n          \"semantic_type\": \"\",\n
\"description\": \"\"\n          }\n          },\n          {\n          \"column\":\n
\"x12\",\n          \"properties\": {\n          \"dtype\": \"number\",\n
\"std\": 77897.45757682854,\n          \"min\": 0.43560000000000004,\n
\"max\": 299077.7344,\n          \"num_unique_values\": 100,\n
\"samples\": [\n          130964.3721,\n          2316.4969,\n
61083.122500000005\n          ],\n          \"semantic_type\": \"\",\n
\"description\": \"\"\n          }\n          },\n          {\n          \"column\":\n
\"x13\",\n          \"properties\": {\n          \"dtype\": \"number\",\n
\"std\": 12666.827095481807,\n          \"min\": 96.04000000000002,\n
\"max\": 38341.5561,\n          \"num_unique_values\": 82,\n
\"samples\": [\n          28334.988900000004,\n
18403.635599999998,\n          7670.256399999999\n          ],\n
\"semantic_type\": \"\",\n          \"description\": \"\"\n          }\n
n          },\n          {\n          \"column\": \"x14\",\n          \"properties\": {\n
\"dtype\": \"number\",\n          \"std\": 2786.7342151991115,\n
\"min\": 0.0,\n          \"max\": 10000.0,\n
\"num_unique_values\": 51,\n          \"samples\": [\n          2.0736,\n
n          23.4256,\n          0.4096\n          ],\n
\"semantic_type\": \"\",\n          \"description\": \"\"\n          }\n
n          }\n          ]\n          }\", \"type\": \"dataframe\", \"variable_name\": \"X\"}

```

## Task 3

Fit OLS on the selected and transformed features and check if the loss has reduced from the baseline estimation.

```

X['y']=y
X

{"summary": "{\n  \"name\": \"X\",\n  \"rows\": 101,\n  \"fields\": [\n
{\n    \"column\": \"x2\",\n    \"properties\": {\n
\"dtype\": \"number\",\n    \"std\": 292.8501773932321,\n
\"min\": -466.86,\n    \"max\": 546.88,\n
\"num_unique_values\": 101,\n    \"samples\": [\n
361.89,\n    73.03,\n    193.86\n    ],\n
\"semantic_type\": \"\",\n    \"description\": \"\"\n    }\n
n    },\n    {\n    \"column\": \"x3\",\n    \"properties\": {\n
\"dtype\": \"number\",\n    \"std\": 55.81221280382463,\n

```

```

{"min": 9.8, "max": 195.81, "num_unique_values": 82, "samples": [168.33, 135.66, 87.58], "semantic_type": "", "description": "", "column": "x4", "properties": {"dtype": "number", "std": 4.942089001927251, "min": 86.83, "max": 108.85, "num_unique_values": 100, "samples": [107.59, 95.67, 104.6], "semantic_type": "", "description": "", "column": "x5", "properties": {"dtype": "number", "std": 30.557704102239096, "min": 0.0, "max": 100.0, "num_unique_values": 51, "samples": [1.44, 4.84, 0.64], "semantic_type": "", "description": "", "column": "x6", "properties": {"dtype": "number", "std": 19555.696239116514, "min": -46686.0, "max": 54688.0, "num_unique_values": 101, "samples": [16733.7936, 143.1388, 1985.1264], "semantic_type": "", "description": "", "column": "x7", "properties": {"dtype": "number", "std": 28863.650047070958, "min": -43149.2355, "max": 59516.9504, "num_unique_values": 101, "samples": [38935.7451, 7295.697, 19831.878], "semantic_type": "", "description": "", "column": "x8", "properties": {"dtype": "number", "std": 37949.18885334589, "min": -86952.675, "max": 96174.59250000001, "num_unique_values": 101, "samples": [22151.2869, 2850.3609, 19791.167400000002], "semantic_type": "", "description": "", "column": "x9", "properties": {"dtype": "number", "std": 5401.921458180213, "min": 976.08, "max": 20034.4095, "num_unique_values": 101, "samples": [6585.5839000000005, 3899.097, 10443.807], "semantic_type": "", "description": "", "column": "x10", "properties": {"dtype": "number", "std": 4751.194534772334, "min": 0.0, "max": 18625.0, "num_unique_values": 101, "samples": [2830.3504000000003, 76.4988, 1045.4016000000001], "semantic_type": "", "description": "", "column": "x11", "properties": {"dtype": "number", "std": 3003.8945434339703, "min": 0.0, "max": 10883.0, "num_unique_values": 101, "samples": [

```



```

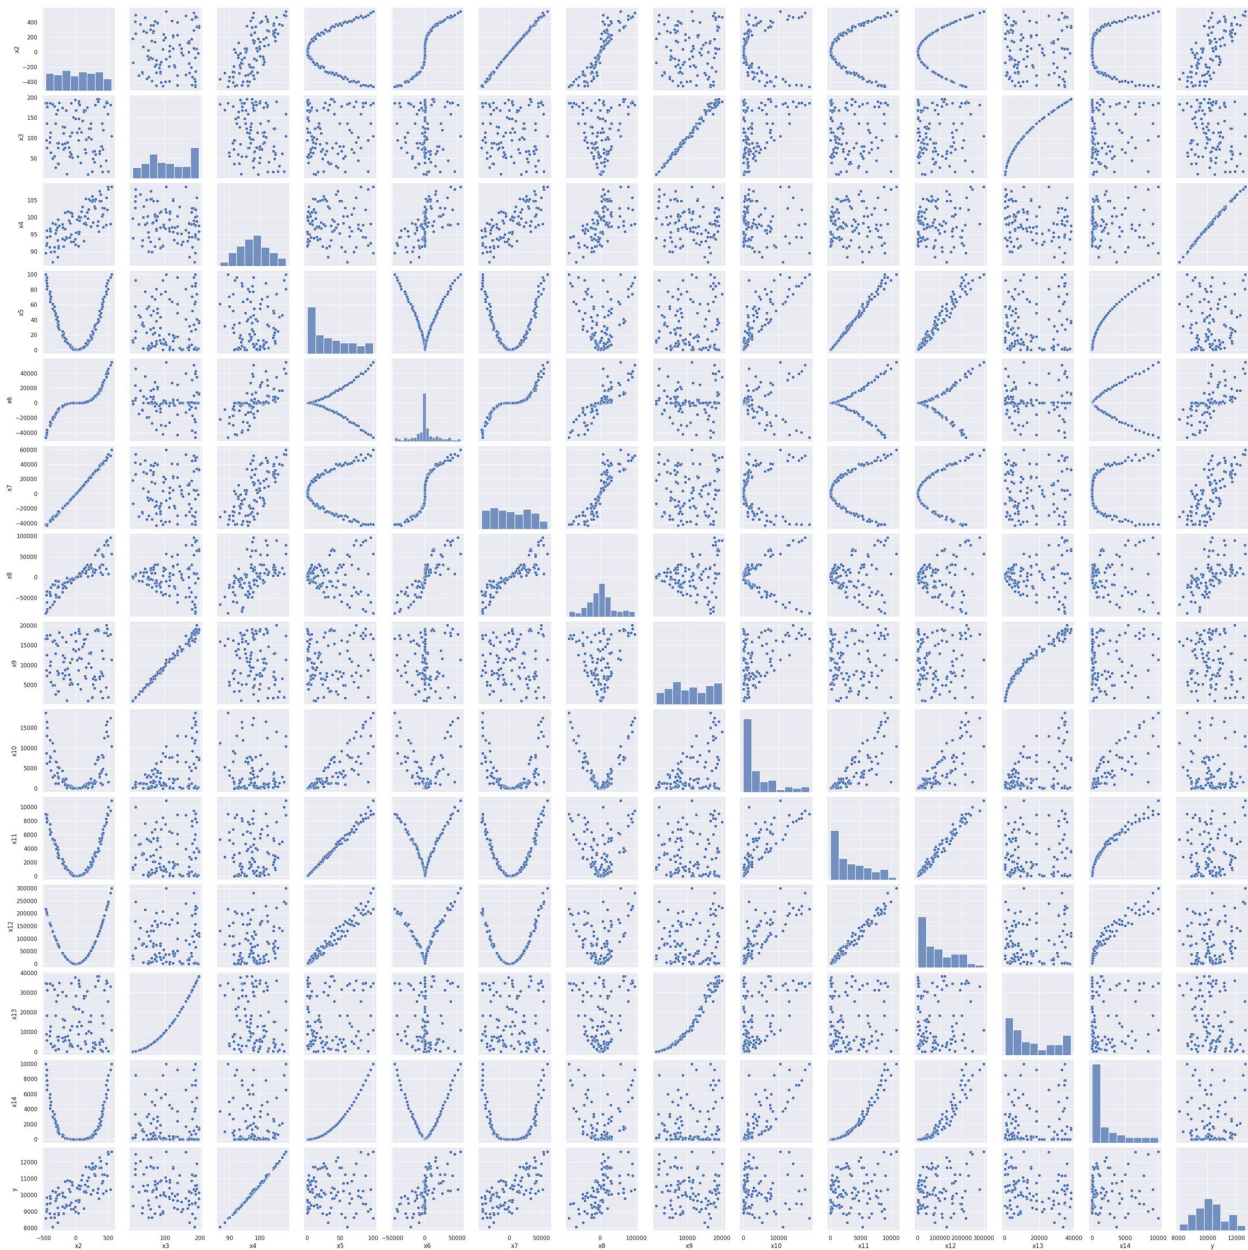
n          4974.9616000000005,\n          195.804,\n1047.552\n    ],\n    \"semantic_type\": \"\",\n\"description\": \"\",\n    },\n    {\n    \"column\":\n    \"x12\",\n    \"properties\": {\n    \"dtype\": \"number\",\n    \"std\": 77897.45757682854,\n    \"min\": 0.43560000000000004,\n    \"max\": 299077.7344,\n    \"num_unique_values\": 100,\n    \"samples\": [\n    130964.3721,\n    2316.4969,\n    61083.122500000005\n    ],\n    \"semantic_type\": \"\",\n    \"description\": \"\",\n    },\n    {\n    \"column\":\n    \"x13\",\n    \"properties\": {\n    \"dtype\": \"number\",\n    \"std\": 12666.827095481807,\n    \"min\": 96.04000000000002,\n    \"max\": 38341.5561,\n    \"num_unique_values\": 82,\n    \"samples\": [\n    28334.988900000004,\n    18403.635599999998,\n    7670.256399999999\n    ],\n    \"semantic_type\": \"\",\n    \"description\": \"\",\n    },\n    {\n    \"column\": \"x14\",\n    \"properties\": {\n    \"dtype\": \"number\",\n    \"std\": 2786.734215199115,\n    \"min\": 0.0,\n    \"max\": 10000.0,\n    \"num_unique_values\": 51,\n    \"samples\": [\n    2.0736,\n    23.4256,\n    0.4096\n    ],\n    \"semantic_type\": \"\",\n    \"description\": \"\",\n    },\n    {\n    \"column\": \"y\",\n    \"properties\": {\n    \"dtype\": \"number\",\n    \"std\": 1022.7661225514418,\n    \"min\": 8062.54,\n    \"max\": 12631.05,\n    \"num_unique_values\": 101,\n    \"samples\": [\n    12266.88,\n    10543.83,\n    11065.37\n    ],\n    \"semantic_type\": \"\",\n    \"description\": \"\",\n    }\n    }\n  ],\n  \"type\": \"dataframe\", \"variable_name\": \"X\"}

```

Visualization of the relationship of the dataset using pairplot.

```
sns.pairplot(X)
```

```
<seaborn.axisgrid.PairGrid at 0x7d66a98a6b60>
```



```
X=X.drop(columns=['y'])
X=sm.add_constant(X)
model=sm.OLS(y,X).fit()
print(model.summary())
```

### OLS Regression Results

```
=====
Dep. Variable:          y    R-squared:
1.000
Model:                  OLS    Adj. R-squared:
```

```

1.000
Method: Least Squares F-statistic:
7.884e+04
Date: Sun, 25 Aug 2024 Prob (F-statistic):
4.23e-171
Time: 10:22:14 Log-Likelihood:
-369.37
No. Observations: 101 AIC:
766.7
Df Residuals: 87 BIC:
803.3
Df Model: 13

```

Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
-----						
const	-9866.6023	76.374	-129.188	0.000	-1e+04	-9714.800
x2	-2.3219	0.108	-21.505	0.000	-2.537	-2.107
x3	3.3032	0.589	5.606	0.000	2.132	4.474
x4	204.2625	0.769	265.534	0.000	202.733	205.791
x5	1.5452	1.015	1.522	0.132	-0.473	3.563
x6	0.0003	0.000	1.606	0.112	-6.98e-05	0.001
x7	0.0242	0.001	22.751	0.000	0.022	0.026
x8	0.0004	9.71e-05	3.877	0.000	0.000	0.001
x9	-0.0320	0.006	-5.549	0.000	-0.043	-0.021
x10	-5.345e-05	0.001	-0.081	0.935	-0.001	0.001
x11	-0.0073	0.011	-0.691	0.492	-0.028	0.014
x12	5.332e-05	0.000	0.412	0.681	-0.000	0.000
x13	-0.0009	0.000	-2.118	0.037	-0.002	-5.3e-05
x14	-7.719e-05	0.001	-0.058	0.954	-0.003	0.003

```
=====
=====
Omnibus:                13.521    Durbin-Watson:
2.402
Prob(Omnibus):          0.001    Jarque-Bera (JB):
14.594
Skew:                   0.893    Prob(JB):
0.000678
Kurtosis:               3.529    Cond. No.
8.90e+06
=====
=====
```

#### Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 8.9e+06. This might indicate that there are strong multicollinearity or other numerical problems.

Sum of squared Errors(SSE) and Root Mean Squared Error(RMSE)

```
predictions=model.predict(X)
residuals=y-predictions
sse2=np.sum(residuals**2)
rmse2=np.sqrt(sse2/101)
print("Sum of Squared Errors(SSE):",sse2)
print("Root Mean Squared Error(RMSE):",rmse2)
```

```
Sum of Squared Errors(SSE): 8879.028866888833
Root Mean Squared Error(RMSE): 9.376096037258227
```

#### Observations 5:

1. SSE is reduced.
2. RMSE is reduced

## Task 4

Install 'lazypredict' package and use the LazyRegressor class to build the regression models. Compare the RMSE reported by all the regression models from LazyRegressor against your OLS losses. Infer the reasons for why different techniques report different performance metrics.

```
from lazypredict.Supervised import LazyRegressor
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=355)
```

```
# Initialize LazyRegressor
```

```
regressor = LazyRegressor(verbose=0, ignore_warnings=True,  
custom_metric=None)
```

```
# Fit the models on the training data and evaluate them on the test data
```

```
models, predictions = regressor.fit(x_train, x_test, y_train, y_test)
```

```
# Display the performance of the models
```

```
print(models)
```

```
100%|██████████| 42/42 [00:02<00:00, 18.18it/s]
```

```
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead  
of testing was 0.000062 seconds.
```

```
You can set `force_col_wise=true` to remove the overhead.
```

```
[LightGBM] [Info] Total Bins 357
```

```
[LightGBM] [Info] Number of data points in the train set: 80, number  
of used features: 13
```

```
[LightGBM] [Info] Start training from score 10290.438037
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain:  
-inf
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain:  
-inf
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain:  
-inf
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain:  
-inf
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain:  
-inf
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain:  
-inf
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain:  
-inf
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain:  
-inf
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain:  
-inf
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain:  
-inf
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain:  
-inf
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain:  
-inf
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain:  
-inf
```

[illegible]

[illegible]

[illegible]



```

-inf
[LightGBM] [Warning] No further splits with positive gain, best gain:
-inf
[LightGBM] [Warning] No further splits with positive gain, best gain:
-inf
[LightGBM] [Warning] No further splits with positive gain, best gain:
-inf
[LightGBM] [Warning] No further splits with positive gain, best gain:
-inf
[LightGBM] [Warning] No further splits with positive gain, best gain:
-inf
[LightGBM] [Warning] No further splits with positive gain, best gain:
-inf
[LightGBM] [Warning] No further splits with positive gain, best gain:
-inf
[LightGBM] [Warning] No further splits with positive gain, best gain:
-inf
[LightGBM] [Warning] No further splits with positive gain, best gain:
-inf
[LightGBM] [Warning] No further splits with positive gain, best gain:
-inf
[LightGBM] [Warning] No further splits with positive gain, best gain:
-inf
[LightGBM] [Warning] No further splits with positive gain, best gain:
-inf
[LightGBM] [Warning] No further splits with positive gain, best gain:
-inf

```

	Adjusted R-Squared	R-Squared	RMSE
\			
Model			
PoissonRegressor	1.00	1.00	12.00
HuberRegressor	1.00	1.00	12.85
BayesianRidge	1.00	1.00	13.50
RANSACRegressor	1.00	1.00	13.57
TransformedTargetRegressor	1.00	1.00	13.57
LinearRegression	1.00	1.00	13.57
LassoLarsIC	1.00	1.00	13.57
LassoLarsCV	1.00	1.00	13.57
RidgeCV	1.00	1.00	18.70
Lasso	1.00	1.00	21.10

LassoCV	1.00	1.00	21.14
LassoLars	1.00	1.00	21.21
SGDRegressor	1.00	1.00	22.40
OrthogonalMatchingPursuitCV	1.00	1.00	22.57
LarsCV	1.00	1.00	22.68
PassiveAggressiveRegressor	1.00	1.00	26.75
Ridge	1.00	1.00	28.32
OrthogonalMatchingPursuit	0.99	1.00	47.92
ExtraTreesRegressor	0.99	1.00	53.18
GradientBoostingRegressor	0.99	1.00	55.83
RandomForestRegressor	0.97	0.99	89.73
DecisionTreeRegressor	0.97	0.99	93.99
BaggingRegressor	0.96	0.99	109.43
XGBRegressor	0.95	0.99	114.14
AdaBoostRegressor	0.94	0.98	125.94
Lars	0.92	0.98	154.00
ExtraTreeRegressor	0.86	0.96	197.57
HistGradientBoostingRegressor	0.78	0.94	249.17
LGBMRegressor	0.78	0.93	251.82
ElasticNet	0.57	0.87	352.88
KNeighborsRegressor	0.35	0.81	431.39
GammaRegressor	0.23	0.77	469.24
TweedieRegressor	0.23	0.77	471.36
ElasticNetCV	0.21	0.76	475.19
NuSVR	-2.37	-0.01	984.03
SVR	-2.41	-0.02	989.50

DummyRegressor	-2.50	-0.05	1003.37
GaussianProcessRegressor	-19.69	-5.21	2438.32
LinearSVR	-349.59	-104.18	10037.14
MLPRegressor	-353.87	-105.46	10098.19
KernelRidge	-367.46	-109.54	10289.68

	Time Taken
Model	
PoissonRegressor	0.04
HuberRegressor	0.04
BayesianRidge	0.01
RANSACRegressor	0.02
TransformedTargetRegressor	0.01
LinearRegression	0.01
LassoLarsIC	0.02
LassoLarsCV	0.04
RidgeCV	0.01
Lasso	0.01
LassoCV	0.08
LassoLars	0.01
SGDRegressor	0.03
OrthogonalMatchingPursuitCV	0.02
LarsCV	0.02
PassiveAggressiveRegressor	0.02
Ridge	0.01
OrthogonalMatchingPursuit	0.02
ExtraTreesRegressor	0.13
GradientBoostingRegressor	0.13
RandomForestRegressor	0.28
DecisionTreeRegressor	0.01
BaggingRegressor	0.04
XGBRegressor	0.69
AdaBoostRegressor	0.10
Lars	0.01
ExtraTreeRegressor	0.02
HistGradientBoostingRegressor	0.05
LGBMRegressor	0.07
ElasticNet	0.01
KNeighborsRegressor	0.01
GammaRegressor	0.01
TweedieRegressor	0.02
ElasticNetCV	0.08
NuSVR	0.02
SVR	0.02

DummyRegressor	0.01
GaussianProcessRegressor	0.02
LinearSVR	0.01
MLPRegressor	0.12
KernelRidge	0.01

#### **Observations 6:**

1. RMSE reported by all the regression models using LazyRegressor is in between 12 to 10289.
2. RMSE reported by the task 1 is 26.67698999975848.
3. RMSE reported by the task 3 after performing EDA is 9.376096037258227

**THE END**